

GERÊNCIA DE INFRAESTRUTURA PARA BIG DATA

Prof. Tiago Ferreto – tiago.ferreto@puccs.br



HADOOP

Big Data and Hadoop History

- Early 2000s – emergence of storage and processing methodologies ("Big Data") from search engine providers – principally Google and Yahoo!
- Search engine providers
 - First group to face with Internet scale problems – how to process and store indexes of all of the documents in the Internet
- Timeline
 - 2003 – Google releases "The Google File System" whitepaper – scalable and fault tolerant file system
 - 2004 – Google releases "MapReduce: Simplified Data Processing on Large Clusters" – large scale processing
 - At the same time – Doug Cutting starts working on a web indexing project called Nutch
 - Inspired by Google papers, he decides to incorporate the storage and processing principles, which are later moved to a new project called Apache Hadoop
 - 2006 – Hadoop is born as an open source project under the Apache Software Foundation

Big Data and Hadoop History

- At the same time as the Hadoop project was created, other technology innovations were happening (data deluge)
 - Rapid expansion of ecommerce
 - Birth and rapid growth of the mobile Internet
 - Blogs and user-drive web content
 - Social media
- Besides Hadoop, it also led to the emergence of other projects, such as Spark, Messaging Systems, NoSQL, etc

What is Hadoop?


- Hadoop is a **data storage and processing platform**
- Central concept → **data locality**
 - Data locality refers to the processing of data where it resides (computation is sent to the data, rather than computation requesting data from its location - e.g. DBMS)
- Big data makes it really difficult to move data around
- Hadoop enables processing large datasets locally on the nodes of a cluster
 - Uses a **shared nothing approach**
 - Each node independently processes a subset of the entire dataset without needing to communicate with one another

What is Hadoop?

- Hadoop is schemaless with respect to its write operations
 - It is a **schema-on-read system**
 - Can store and process a wide range of data, such as:
 - Unstructured text documents
 - Semi-structured JSON or XML documents
 - Well structured extracts from relational database systems
- In **schema-on-write** systems, such as relational databases, data is strongly typed and a schema is predefined and enforced in all operations
 - NoSQL databases are **schema-on-read** systems
- Pros: doesn't need indexes, statistics or optimization constructs
- Cons: datasets with larger size and replicated data
- Hadoop uses a **divide and conquer approach** to tackle problems applying data locality and shared nothing concepts

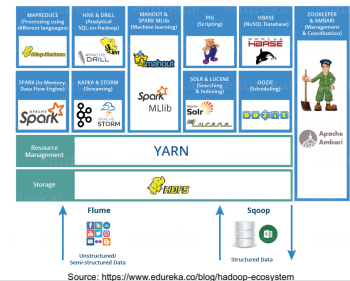
Hadoop –Components

- Hadoop Common – libraries and utilities needed by other modules
- HDFS (Hadoop Distributed File System) – distributed file-system for data storage
- Hadoop YARN (Yet Another Resource Negotiator) – platform for managing computing resources and scheduling applications
- Hadoop MapReduce – implementation of the MapReduce programming model



Hadoop Ecosystem Components

- Hadoop Ecosystem includes any other projects that interact or integrated with Hadoop in some way (Data ingestion, Data analysis, NoSQL Databases, Security, Management, Search, Machine Learning, Workflow coordination, etc)



Commercial Hadoop Landscape

- 2008 – Cloudera
 - CDH (Cloudera Distribution of Hadoop)
- 2009 – MapR
 - Includes MapRFS – custom storage solution, adapted from HDFS
- 2011 – HortonWorks
 - HDP (Hortonworks Data Platform)
- Cloudera, Hortonworks, and MapR are considered “pure play” Hadoop vendors
 - Business models are founded upon Hadoop
 - Other vendors with proprietary distributions – IBM, Pivotal, Teradata
 - Hadoop is not their core business
- 2014 – ODPI (Open Data Platform Initiative) founded by Hortonworks – odpi.org
 - Goal: provide a consistent, commercial ready set of Hadoop core and selected ecosystem components
 - Includes other vendors: IBM, Teradata, EMC², SAS, Pivotal, Splunk, and others

Hadoop releases

- Hadoop 1 (2012)
 - First release - HDFS and MapReduce
- Hadoop 2 (2013)
 - Introduces YARN - optimizes resource allocation for various applications
- Hadoop 3 (2017)
 - Support multiple name nodes on HDFS (solves SPOF problem)
 - HDFS uses erasure coding – reduces storage overhead
 - Utilization of containers (Docker)
 - Supports the utilization of GPUs for processing (Deep Learning)


Typical Hadoop Use Cases

- **Data warehouse offload**
 - Utilization of Hadoop for long-term storage and ETL routines
 - Leverages the lower cost storage and distributing processing capabilities of Hadoop (in comparison to specialized Data Warehouse platforms)
 - Implementation of a Data Lake - storage of unstructured, unprocessed data prior to being staged of integrated into a data warehouse or exposed in a data mart

Typical Hadoop Use Cases

- **Event and complex event processing**
 - Associated with IoT (Internet of Things) - thousands of end points, including temperature sensors, RFID, NFC scanners, CCTV cameras, signalling systems, etc.
 - Involves the ingestion and processing of streaming data sources, such as sensor data, message data, or log data
 - Hadoop provides low-cost storage to accumulate millions of messages and a computational platform to mine volumes of messages to identify patterns or peculiarities.
 - Typically involves other Hadoop ecosystem platforms, such as Storm, Flume, Spark Streaming, Kafka, or others

HDFS Cluster Overview



YARN Basic Concepts

- Uses a master-slave model like HDFS
 - ResourceManager and NodeManager
- ResourceManager
 - Master of the YARN Cluster
 - Responsible for scheduling applications and granting cluster compute resources
 - Tracks applications' status
 - Resources are represented as units called "containers" (combinations of CPU cores and memory)
 - Tracks available capacity in the cluster - allocates and releases containers to applications
 - For each application submitted, ResourceManager allocates the first container on an available NodeManager as a delegate process for the application called "ApplicationMaster"
 - ApplicationMaster negotiates further containers required for the application
 - ResourceManager serves a web interface on port 8088


ResourceManager Web Interface

The screenshot shows the 'All Applications' page of the ResourceManager web interface. It displays a table with columns for Application ID, Name, User, Application Type, Queue, Status, and Progress. The table lists several applications, including 'hadoop-job-1' and 'hadoop-job-2', with their respective statuses and progress bars.

YARN Basic Concepts


- NodeManager
 - Slave node daemon
 - Manages containers on the host
- ApplicationMaster
 - First container allocated by the ResourceManager to run on a NodeManager for an application
 - Determines resources required to run the application (based on the amount of data to be processed)
 - Orchestrates the allocation of resources to execute application stages (eg. Map and reduce stages)
 - ApplicationMaster requests resources from the ResourceManager on behalf of the application
 - ResourceManager grants resources on the same or other NodeManagers to the ApplicationMaster
 - Monitors progress of tasks, stages (groups of tasks executed in parallel), and dependencies

YARN Cluster Overview



Hadoop Cluster Architecture

- Combines HDFS and YARN Clusters



Hadoop Cluster Deployment Modes

- Hadoop supports three deployment modes
 - LocalJobRunner (Standalone Operation)
 - Runs an application locally in a single JVM
 - Enables integrating the application with an IDE (e.g. Eclipse) for performing unit tests, debugging, and tracing
 - Uses the local filesystem instead of HDFS
 - Pseudo-Distributed
 - All Hadoop daemons are executed on separate JVMs on a single host
 - Simulates how a fully cluster would function
 - Useful for testing to simulate the interaction between components in a real cluster using a single machine
 - Fully Distributed Cluster
 - Executes master and slave daemons on distinct machines
 - Typical deployment mode for production systems

Pseudo-Distributed Mode



Hadoop Deployment

- Hadoop was originally developed for Linux
 - Windows-compatible distribution already exists, but may present compatibility issues with other Hadoop ecosystem projects
- Recommendations
 - Do not use LVM (Logical Volume Manager). It may restrict performance (especially on slave nodes)

Requirements (medium to large scale production Hadoop clusters)


- Master Nodes**
 - 16 or more CPU cores (preferably 32)
 - 128GB or more RAM (preferably 256GB)
 - RAID Hard Drive Configuration (preferably with hot-swappable drives) - for fault tolerance
 - Redundant power supplies - for fault tolerance
 - Bonded Gigabit Ethernet or 10Gigabit Ethernet

Requirements (medium to large scale production Hadoop clusters)

- Slave Nodes**
 - 16-32 CPU cores
 - 64-512 GB of RAM
 - 12-24 1-4 TB hard disks in a JBOD Configuration
 - JBOD (Just a Bunch of Disks) – directly attached storage not in RAID configuration; each disk operates independently
 - RAID is not recommended for block storage on slave nodes!
 - Access speed is limited by the slowest disk in the array. JBOD has been proven to outperform RAID 0 for block storage by 30% to 50% in benchmarks conducted at Yahoo!
- Slave nodes support failures, therefore they do not require the same degree of fault tolerance that master nodes do
- Storage capacity in slave nodes should not be fully allocated!
 - Failure of a node may impact the network due to replication of blocks

Networking Considerations

- Fully-distributed implementations consume significant amount of network resources
 - Control messages, status updates and heartbeats, block reports, data shuffling, and block replication
- On-premises deployment
 - Private subnets with dedicated switches
 - Redundant core and TOR (Top of Rack) switches
 - Utilization of STP (Spanning Tree Protocol) to avoid loops
- Hostname resolution between nodes
 - Forward and reverse DNS configuration or hosts file
- Time synchronization (NTP – Network Time Protocol)
 - Used by some components (e.g. Kerberos)



Gerencia de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

Software requirements

- Hadoop is almost entirely written in Java
- Requirements
 - Java Runtime Environment (JRE) 1.7 or above
 - Java Development Kit (JDK) 1.7 or above - for compiling MapReduce applications
- OpenJDK and Oracle versions supported
- Some versions may present problems!
- Site with tested JDKs
 - <https://wiki.apache.org/hadoop/HadoopJavaVersions>
- Other ecosystem projects may have other specific requirements
 - For instance, Apache Spark requires Scala and Python

Gerencia de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

Installing Hadoop On-premises

- Non-Commercial Hadoop
 - Install from Apache builds – <http://hadoop.apache.org>
- Commercial Distributions
 - Cloudera (<http://www.cloudera.com/>)
 - Provides a Quickstart VM - <http://www.cloudera.com/downloads.html>
 - Cloudera Manager (proprietary software) – provides a management console and framework to deploy, manage and monitor Hadoop clusters
 - Hortonworks (<http://hortonworks.com/>)
 - Hortonworks Sandbox – VM with Hadoop pseudo-distributed mode - <http://hortonworks.com/products/sandbox/>
 - Hortonworks Data Platform (distribution with complete Hadoop stack and selected ecosystem components)
 - Uses Apache Ambari for deployment and management
 - MapR (<https://www.mapr.com/>)
 - MapR Sandbox – VM with Hadoop pseudo-distributed mode - <https://www.mapr.com/products/mapr-sandbox-hadoop/download>
 - MapR Control System (MCS) – central console to configure, monitor and manage MapR clusters

Gerencia de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

Installing Hadoop on the Cloud

- Amazon Web Services
 - EMR (Elastic MapReduce) – Amazon's Hadoop-as-a-Service platform
 - Enables creating clusters with a specific number of nodes, node instance types (EC2), Hadoop distribution, and additional ecosystem applications
 - Uses S3 for reading/writing data
 - Provisioned on demand
 - <https://aws.amazon.com/emr/>



Gerencia de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

Installing Hadoop on the Cloud

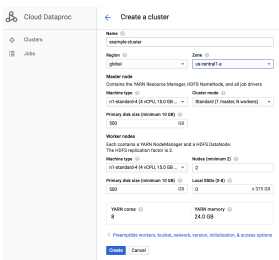
- Microsoft Azure HDInsight
 - <https://azure.microsoft.com/en-us/services/hdinsight/>



Gerencia de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

Installing Hadoop on the Cloud

- Google Cloud Dataproc
 - <https://cloud.google.com/dataproc/>



APACHE HADOOP - HANDS-ON