

GERÊNCIA DE INFRAESTRUTURA PARA BIG DATA

Prof. Tiago Ferreto – tiago.ferreto@puccs.br



MAPREDUCE


Introduction to MapReduce

- Whitepaper "MapReduce: Simplified Data Processing on Large Clusters" released in December 2004 from Google
 - High-level description of Google's approach to processing, specifically indexing and ranking, large volumes of text data for search-engine processing
- Influence on the Nutch project (Yahoo!)
 - Creators of the Nutch project (including Doug Cutting), incorporated the principles outlined in the Google MapReduce and Google File System papers into the project now known as Hadoop

Motivation

- Limitations on the scale-up approach to increase processing capacity
- Before MapReduce (2004) there were several programming frameworks for distributed systems – Message Passing Interface (MPI), Parallel Virtual Machine (PVM), HTCondor, and others.
- But they had several limitations:
 - **Complexity in programming:** need to explicitly handle state and synchronization between distributed processes, including temporal dependencies
 - **Partial failures (difficult to recover from):** synchronization and data exchange between processes in a distributed system made dealing with partial failures much more challenging
 - **Bottlenecks in getting data to the processor:** most distributed systems sourced data from shared or remote storage
 - **Limited scalability:** finite bandwidth between processes limit how the distributed systems can scale

Scalability issues in distributed systems



Design goals for MapReduce

- **Automatic parallelization and distribution:** the programming model should make it easy to parallelize and distribute computations
- **Fault tolerance:** the system must be able to handle partial failure. If a node or process fails, its workload should be assumed by other functioning components in the system.
- **Input/output (I/O) scheduling:** Task scheduling and allocation aim to limit the amount of network bandwidth used. Tasks are dynamically scheduled on available workers so that faster workers process more tasks.
- **Status and monitoring:** Status of each component and its running tasks, including progress and counters, are reported to a master process. This makes it easy to diagnose issues, optimize jobs, or perform system capacity planning.

MapReduce – Key-Value Pairs and Records

- Input, output, and intermediate records in MapReduce are represented in the form of key-value pairs
 - Key-value pairs are commonly used in programming to define a unit of data
- Key** is usually an identifier (e.g., name of the attribute)
 - In some systems, the key needs to be unique with respect to other keys in the same system (e.g., NoSQL key-value stores), but this is NOT REQUIRED in MapReduce
- Value** is the data that corresponds to the key
 - It can be a simple, scalar value such as an integer, or a complex object such as a list of other objects

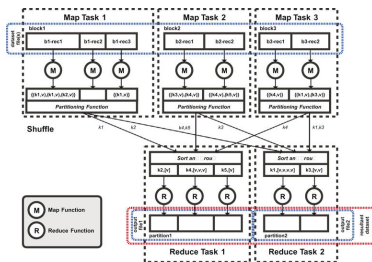
Key	Value
City	Chicago
Temperatures	[35,38,27,16]

MapReduce – Key-Value Pairs and Records

- Key-value pairs are implemented in many programming languages
 - Python uses dictionaries. Ruby uses hashes.
- Key-value pairs are the atomic data unit used for processing in MapReduce programming
- Complex problems are often decomposed in Hadoop into a series of operations against key-value pairs

MapReduce Programming Model

- Inspired by the map and reduce primitives in Lisp and other functional programming languages
- MapReduce includes two developer-implemented processing phases, the **Map phase** and the **Reduce phase**, along with a **Shuffle-and-Sort phase**, which is implemented by the processing framework




Map Phase

- Initial phase for processing an input dataset
- It uses input format and record reader functions to derive records in the form of key-value pairs for the input data
- Map phase applies a function or functions to each key-value pair over a portion of the dataset
 - When using HDFS, it is block in the filesystem
 - n blocks of data in the input dataset → n Map tasks (also referred to as Mappers)
 - More Map tasks may be generated due to node failures
- No state is shared between processes
- Each Map task iterates through its portion of the dataset in parallel with the other Map tasks
- Every record (key-value pair) is processed once and only once (with exceptions only for task failure or speculative execution)

Map Phase

- Example
 - Three Map tasks operating against three filesystem blocks (block1, block2, and block3)
 - Each Map task calls its map() function (M) once for each record (key-value pair)
 - Each map() function receives one key-value pair and outputs zero or more key-value pairs
 - Results are considered intermediate data (may be further processed by the Reduce phase)



Map Phase

- Pseudo-code
 - `map (in_key, in_value) → list (intermediate_key, intermediate_value)`
- Examples
 - Filtering log messages (only ERROR messages)
 - `let map (k, v) = if (ERROR in v) then emit (k, v)`
 - Manipulate values (convert to lowercase)
 - `let map (k, v) = emit (k, v.toLowerCase())`

Map Phase

- Any map() function is valid as long as the function can be executed against a record contained in a portion of the dataset, in isolation from other Map tasks in the application that are processing other portions of the dataset
 - There may be no dependencies between Map tasks
- Map task collects lists of intermediate data key-value pairs emitted from each map function into a single list grouped by the intermediate key
- Combined list of intermediate values grouped by their intermediate keys is then passed to a **partitioning function**

Partitioning Function (or Partitioner)

- Goal: ensure each key and its list of values is passed to one and only one Reduce task or Reducer
 - Most common implementation: hash partitioner
 - Creates a hash (or unique signature) for the key and divides the hashed key space into n partitions (where n is the number of Reducers)
- Custom partitioners can also be implemented
 - Example: implement a Partitioner to partition by the month in order to process a year's worth of data
- Partitioning function is called for each key with an output representing the target Reducer for the key, typically a number between 0 and $n - 1$ (n = number of Reducers)

Shuffle and Sort

- The output from each separate Map task is sent to a target Reduce task as specified by the application's partitioning function
- Requires data to be physically transferred between nodes, requiring network I/O and consuming bandwidth
- Keys and their values are grouped together and presented in key-sorted order to the target Reducer (SORT)
 - For instance, if the key is a Text value then the keys would be presented to the Reducer in ascending alphabetical order

Reduce Phase

- Only starts when
 - all of the Map tasks have completed AND
 - Shuffle phase has transferred all of the intermediate keys and their lists of intermediate values to their target Reducer (or Reduce task)
- Each Reduce task (or Reducer) executes a reduce() function for each intermediate key and its list of associated intermediate values
- The output from each reduce() function is zero or more key-value pairs considered to be part of the final output
 - Output may be the input to another Map phase in a complex multistage computational workflow
 - In the context of the individual MapReduce application, the output from the Reduce task is final

Reduce Phase

- Pseudo-code representation:
 - reduce (key, list (intermediate_value)) → key, out_value
- reduce() functions are often aggregate functions such as sums, counts, and averages
- Example: Sum Reducer

```
let reduce (k, list <v>) =
  sum = 0
  for int i in list <v> :
    sum + = i
  emit (k, sum)
```


Fault Tolerance

- MapReduce was designed to tolerate node failures
- If a Map task fails, it will automatically be rescheduled by the master process on another node, preferably a node that has a copy of the same block(s), maintaining data locality
- A task can fail and be rescheduled **four times** before the job is deemed to have failed
- If a Reduce task fails, it also can be rescheduled and its input data resupplied
 - intermediate data is retained for the life of the job

Combiner Functions

- In the case of commutative and associative Reduce operations (e.g., sums and counts), operations can be performed after the Map task is complete on the node executing it (before the Shuffle-and-Sort phase)
 - Utilization of a Combiner function, or Combiner
- Non commutative and associate operations (e.g., averages) cannot be implemented as a combiner
 - $\text{Avg}(\text{LIST}) \neq \text{Avg}(\text{Avg}(\text{SUBLIST}), \text{Avg}(\text{SUBLIST}))$
 - Example: $\text{Avg}(2, 2, 3, 3, 3) \rightarrow 2.6$, while $\text{Avg}(\text{Avg}(2, 2), \text{Avg}(3, 3, 3)) = \text{Avg}(2, 3) \rightarrow 2.5$
- Using a combiner reduces the amount of data transferred in the Shuffle phase and reduce the computational load in the Reduce phase
- Combiner function is often equal as the reduce() function, but executed on the Map task node

Combiner function




Asymmetry and Speculative Execution

- Map and Reduce phases are asymmetrical
 - One instance of a Map task may do more processing than other Map tasks mapping over the same dataset
 - Example: filtering weblog for a specific IP \rightarrow some blocks may contain more references to the IP than others
- Some Map tasks may run slower than others
 - Map phase must complete before Reduce phase starts \rightarrow performance problem: need to wait slower mappers
- Speculative execution (Governed by the ResourceManager and ApplicationMaster)
 - Looks for configurable, tolerable difference in progress between tasks
 - If a task falls outside this tolerance (is taking too long to complete), a duplicate task is created to process the same data
 - The results of the first task to complete are used and the other task is killed (and output discarded)
 - Prevents a slow, overloaded, or unstable node from becoming a bottleneck

Map-only MapReduce Applications

- Map-only MapReduce application \rightarrow a MapReduce application with zero Reduce tasks
- Common applications
 - ETL routines where the data is not intended to be summarized, aggregated, or reduced
 - file format conversion jobs
 - image processing jobs
- There is no partitioning function \rightarrow the output from the Map task is considered to be the final output
- Provides massive parallelization avoiding the expensive Shuffle-and-Sort operation



MapReduce implementation in Hadoop

- Two implementations
 - MapReduce cluster framework (MR1 or MapReduce version 1) – Hadoop 1
 - YARN (MR2) – Hadoop 2
- MR1 uses the following daemons
 - JobTracker (instead of the ResourceManager in YARN)
 - TaskTracker (instead of the NodeManager in YARN)
- Drawbacks of MR1
 - Does not work with non-MapReduce programs (e.g., Spark)
 - Limited scalability
 - Inefficient usage of processing capacity (especially regarding heterogeneous resources)

Running and Application on YARN

- YARN schedules and orchestrates applications and tasks in Hadoop
 - Applies data locality concept – tasks are scheduled on the node where data resides
- Application's workload is distributed across NodeManagers
 - NodeManagers are responsible for carrying out tasks
- ResourceManager (YARN's master) is responsible for
 - Assigning an ApplicationMaster (delegate process for managing the execution and status of an application)
 - Keeping track of available resources on the NodeManagers, such as CPU cores and memory
- Compute and memory resources are presented to applications in processing units called **containers**
- ApplicationMaster determines container requirements for the application and negotiates these resources with the ResourceManager

